

AITK & DSDL — Usage Flow Guide

What each tool is for, how the data flows, and worked examples — from in-Splunk ML commands to containerized deep learning.

1. The big picture

Splunk's machine-learning stack is two complementary layers. You almost always start with AITK; you reach for DSDL only when you need heavy or custom models.

AITK Splunk AI Toolkit

Formerly the **Machine Learning Toolkit (MLTK)**.

Runs ML **inside Splunk** using SPL. Ships ~30 classic algorithms (regression, clustering, classification, anomaly detection, forecasting), guided *Assistants*, and the `fit` / `apply` search commands. No containers, no Python to write.

Best for: tabular data, quick models, dashboards, well-known algorithms.

DSDL Data Science & Deep Learning


Formerly the **Deep Learning Toolkit (DLTK)**. **Extends** AITK.

Runs your model in an **external Docker container** (TensorFlow, PyTorch, NLP/LLM libraries) with a **JupyterLab** dev environment. Splunk streams data to the container's API and gets predictions back.

Best for: neural nets, NLP/LLMs, GPU training, custom Python, anything beyond the built-in algorithms.

Relationship: DSDL is built *on top of* AITK. It adds the `MLTKContainer` algorithm to the same `fit` / `apply` SPL you already use in AITK — but the computation happens in the container instead of in Splunk. Install order is always **PSC → AITK → DSDL**.

2. How the data flows



Both expose the **same verbs** — `fit` (train) and `apply` (score). The only difference is *where* the math runs: in Splunk (AITK) or in the container (DSDL, via `algo=MLTKContainer`).

3. AITK vs DSDL — the difference in depth

Clear up the common misconception first. It is *not* "AITK = ready-made models you just use, DSDL = where you train." **Both train models on your own data** using the same `fit` command. The real difference is *how you get the model*: with AITK you **pick an algorithm from a fixed catalog**; with DSDL you **write the model yourself in code**.

The one-line distinction

AITK — choose from ~30 ready-made algorithms and train them on your data; you tune parameters, not the model design. **DSDL** — author any model from scratch in Python (any framework, any architecture) and run it in a container.

AITK = order from the menu

You pick a dish (algorithm) off the menu and the kitchen cooks it with *your* ingredients (data). Fast and easy — but you can only order what's on the menu, and you adjust seasoning (parameters), not the recipe itself.

DSDL = cook in the kitchen

You step into the kitchen and write the recipe yourself — any ingredients, any technique, any equipment (GPU). Total freedom at the cost of having to write and maintain the code and run the container.

Detailed comparison

Aspect	AITK	DSDL
Trains on your data (<code>fit</code> / <code>apply</code>)	Yes	Yes
How you obtain a model	Pick one of ~30 built-in algorithms	Write it yourself in a Jupyter notebook
Algorithm catalog	Fixed (regression, classification, clustering, density/anomaly, forecasting, preprocessing)	Unlimited — whatever you code
Custom model architecture	No	Yes — define layers, loss, training loop
Frameworks under the hood	scikit-learn, StatsModels (fixed)	TensorFlow, PyTorch, HuggingFace, XGBoost, spaCy... anything <code>pip</code> -installable
Deep learning / neural nets	Very limited	Full
NLP / transformers / LLM / embeddings	No	Yes (incl. RAG)
GPU training	No	Yes (GPU golden image)
Custom preprocessing / feature engineering	Within SPL only	Arbitrary Python in the notebook
Where computation runs	Inside Splunk (PSC add-on)	External Docker container
Coding required	None — SPL only	Python (notebook)
Where the model is stored	In Splunk (model lookups)	In the container volume (<code>/srv/app/model</code>)
Setup & operations	Just install the app	Docker + container lifecycle + endpoint config
Latency / overhead	Low (in-process)	Network hop; container must be running
Extensibility	Limited (custom algos possible but constrained)	Effectively unlimited
Sweet spot	Classic ML on tabular/log data, dashboards, quick wins	DL, NLP/LLM, GPU, custom research

What only **DSDL** can do

- Define your own network architecture (CNN / RNN / Transformer)
- Use PyTorch / TensorFlow / HuggingFace models
- Fine-tune or serve LLMs; build embeddings / RAG
- Train on a **GPU**
- Custom loss functions, custom training loops
- Arbitrary Python preprocessing + any **pip** package

...but **AITK** wins when

- A standard algorithm already fits the problem
- You want **zero code** and instant dashboards/Assistants
- You'd rather not run & maintain Docker
- You need low latency / simple operations
- The data is tabular and the model is "classic ML"

Rule of thumb: try AITK first. If a built-in algorithm solves it, you're done with no code. Move to DSDL only when you hit the catalog's limits — e.g. our **DGA detector** is a char-level CNN, which AITK doesn't offer, so it lives in a DSDL notebook.

4. Using AITK — examples

AITK adds `fit` / `apply` plus helper commands. Models are saved in Splunk and reused by name. A few representative flows:

4.1 Forecasting / regression

```
| inputlookup server_power.csv
| fit LinearRegression "ac_power" from "total-cpu-utilization" "total-disk-utilization"
  into power_model
| apply power_model
| eval residual = 'ac_power' - 'predicted(ac_power)'
```

4.2 Anomaly detection (unsupervised)

```
index=web sourcetype=access_combined
| timechart span=5m count AS hits
| fit DensityFunction hits into hits_density
| apply hits_density
| where IsOutlier=1      ← flag unusual traffic spikes
```

4.3 Clustering

```
index=fw | stats sum(bytes) AS bytes count AS sessions by src_ip
| fit KMeans bytes sessions k=4 into ip_clusters
| apply ip_clusters      ← adds a "cluster" field per src_ip
```

4.4 Classification

```
| inputlookup firewall_traffic.csv
| fit RandomForestClassifier used_by_malware from src_bytes dest_bytes duration
  into fw_class
| apply fw_class
```

Assistants: AITK also ships point-and-click *Assistants* (Predict Numeric/Categorical Fields, Detect Outliers, Forecast Time Series, Cluster Events) that generate this SPL for you — a great way to learn the commands.

5. Using DSDL — examples

DSDL routes `fit` / `apply` to a container running your model. You author the model once in JupyterLab; Splunk then calls it like any other algorithm.

5.1 The notebook contract

A DSDL model is a notebook whose cells are tagged with magic comments. On **save**, DSDL compiles the tagged cells into `/srv/app/model/<name>.py`:

Cell tag	Function	Role
# mltkc_import	imports	module top
# mltkc_init	init(df,param)	build the model
# mltkc_fit	fit(model,df,param)	train
# mltkc_apply	apply(model,df,param)	score → DataFrame
# mltkc_save / _load	save() / load()	persist / reload
# mltkc_summary	summary(model)	model info

5.2 The development loop

1. **Stage** real data into the container (no training):

```
| inputlookup data.csv | fit MLTKContainer mode=stage algo=my_model ...
```

2. **Iterate** in JupyterLab: `df,param = stage("my_model")` → run `init/fit/apply` by hand, tweak the cells.

3. **Save** the notebook → the `.py` is rebuilt.

4. **Train & score** from Splunk (no `mode=stage`):

```
... | fit MLTKContainer algo=my_model ... into app:my_model then ... | apply my_model .
```

5.3 Worked example — DGA domain detection on BOTSv1

A character-level neural network classifies algorithmically-generated malware domains vs benign ones, then scores real DNS queries from the BOTSv1 dataset.

```
# 1) TRAIN — stream labeled domains to the container, train + save the model
| inputlookup dga_training_domains.csv
| fit MLTKContainer algo=dga_neural_network epochs=25 is_dga from domain
  into app:dga_model

# 2) SCORE — rank BOTSv1 DNS queries by how "DGA-like" they look
index=botsv1 sourcetype=stream:dns message_type=Query
| eval domain=lower(mvindex('query{}',0))
| stats count by domain
| apply dga_model
| where is_dga_predicted=1
| sort - dga_score
| table domain count dga_score
```

Why DSDL here? A trainable char-level CNN (Keras/TensorFlow) isn't a built-in AITK algorithm — so the model lives in the container, while the SPL stays identical in spirit to AITK's `fit` / `apply`.

6. End-to-end flow (this lab)

1. Install

2. Configure DSDL

3. Container

4. Develop

5. Train

6. Score / detect

PSC → AITK → DSDL (Splunk apps)

Setup page → Docker mode → Test & Save

golden image runs as compose service "mltk-dev" (DEV mode)

:5000 model API :8888 JupyterLab :6006 TensorBoard

open https://localhost:8888 → edit notebook → Save

| fit MLTKContainer algo=... into app:model

| apply model → dashboards, alerts, notable events

Key facts for this environment

Item	Value
Splunk Web	http://localhost:8000 (admin / your password)
JupyterLab	https://localhost:8888 — HTTPS, password splunkdssl
Model API (fit/apply)	https://host.docker.internal:5000 (Endpoint URL in DSDL)
DSDL Docker Host	tcp://docker-proxy:2375 (sidecar; Splunk can't read the raw socket)
Container	compose service mltk-dev — do not also start one in the DSDL UI

Common gotchas.

• JupyterLab is HTTPS — http://localhost:8888 returns "didn't send any data". Use https:// and accept the self-signed cert.

• Container must be running for fit / apply to work.

• Edited a notebook but the search runs old code? You forgot to Save (the .py only regenerates on save).

• Check Hostname = Disabled in DSDL Setup (self-signed dev certs).

7. Quick decision guide

Your situation	Use
"Find anomalies / forecast / cluster my logs"	AITK built-in algorithm or Assistant
"Predict a number/category from tabular fields"	AITK regression / classification
"I need a neural net / transformer / embeddings / LLM"	DSDL custom notebook
"I have Python model code / a specific framework"	DSDL container
"I need a GPU to train"	DSDL (GPU golden image)